



UNIVERSITÀ DI PISA

Computer Engineering

Artificial Intelligence and Data Engineering

Project Documentation of
*Concurrent Database
Access Simulation*

Authors:

Rotelli Luca - Sfar Omar Tomàs - Gashi Greta

Academic Year: 2025-2026

Contents

1	Introduction	1
1.1	Objectives	1
1.2	Context	1
2	Problem Definition	2
2.1	Problem Statement	2
2.2	Constraints	2
2.3	Key Performance Metrics	2
3	System Implementation and Verification	3
3.1	System Model	3
3.1.1	Architecture	3
3.1.2	Temporal Model	3
3.2	Modules and Components	3
3.2.1	User Module	3
3.2.2	Table Module	3
3.2.3	Network	4
3.2.4	Module Parameters	4
3.3	Implementation Details	4
3.3.1	Design Choices	4
3.3.2	Concurrency Management Algorithm	4
3.4	Verification	5
3.4.1	Degeneracy Test	5
3.4.2	Continuity Test	5
3.4.3	Consistency Test	6
3.4.4	Theoretical Verification	7
3.4.5	Throughput Analysis	10
3.5	Conclusions	11
4	Data Analysis and Conclusions	12
4.1	Warm-Up Period Analysis	12
4.2	Experimental Design	13
4.3	Factor Impact Analysis	13
4.3.1	Key Findings	14

4.4	Model Validation	14
4.4.1	Testing Normality Hypothesis	14
4.4.2	Homoskedasticity	15
4.4.3	Residual Magnitude Analysis	16
4.5	Conclusions	17
4.5.1	System Capacity Analysis	18
5	Final Conclusions	20
5.1	Main Results	20
A	Appendix: Detailed Configuration Table	21

1 Introduction

This project implements a discrete event simulation of a concurrent database access system using the **OMNeT++ framework** (Discrete Event Simulator).

1.1 Objectives

The main objective is to evaluate the performance of a multi-client database system by implementing a readers/writers model with concurrency control through mutual exclusion and FCFS (First Come First Served) queueing.

1.2 Context

The simulation analyzes how the system responds to varying loads of read and write operations from N concurrent users accessing M database tables.




Figure 1.1: Simplified interaction schema: N Users accessing M Tables concurrently.

2 Problem Definition

2.1 Problem Statement

- N concurrent users randomly access M database tables.
- Each user generates requests according to a Poisson process (rate λ).
- Each request is a **READ** with probability p , or **WRITE** with probability $(1 - p)$.
- Table selection follows a specified distribution (uniform or lognormal).
- Each access requires a fixed service time S .

2.2 Constraints

- **Reads:** Multiple simultaneous reads on the same table are allowed (reader lock).
- **Writes:** Exclusive access only (write lock), blocked by active readers.
- **Queuing:** FCFS to handle conflicting requests.
- **Mutual Exclusion:** Implemented via reader/writer counters.

2.3 Key Performance Metrics

- Average wait time (from request submission to response).
- Throughput (completed accesses per second).
- Maximum/average queue length per table.
- Table utilization (fraction of time occupied).
- Percentage of completed operations.

3 System Implementation and Verification

3.1 System Model

3.1.1 Architecture

- **Network:** Fully-connected mesh topology with N users and M tables.
- **Communication:** Asynchronous message-passing between modules.
- **Timing:** Discrete event simulation with `simTime()` in seconds.

3.1.2 Temporal Model

- **Inter-arrival times:** Exponential with rate $\lambda = 1/T_{inter}$. Generator: `exponential(1/lambda)`.
- **Table selection:**
 - Uniform: `intuniform(0, numTables-1)`.
 - Lognormal: `lognormal(m, s)` mapped to $[0, M - 1]$.
- **Response time:** Queueing time + Service time.

3.2 Modules and Components

3.2.1 User Module

Defined in `User.h/cc`, `User.ned`. It generates database access requests.

Behavior: Schedules accesses per Poisson process with rate λ . For each access, it selects a table and type (read/write), sends the request, and records statistics upon response.

3.2.2 Table Module

Defined in `Table.h/cc`, `Table.ned`. Simulates concurrent access to a single table.

Internal State: `activeReaders`, `writeActive`, `requestQueue` (FCFS), `serviceEvents`.

Mutual Exclusion Logic: • If `writeActive=true`: all new accesses blocked.

- If `activeReaders>0`: new reads OK, new writes blocked.
- If `activeReaders=0` and `writeActive=false`: both reads and writes OK.

3.2.3 Network

Defined in `DatabaseNetwork.ned`.

- **Parameters:** `numUsers` (default 60), `numTables` (default 20).
- **Topology:** Fully-connected mesh where each `user[i]` sends requests to any `table[j]`.

3.2.4 Module Parameters

Table 3.1 summarizes the configurable parameters:

Module	Parameter	Description
User	<code>lambda</code>	Rate of exponential inter-arrival time (λ)
User	<code>readProbability</code>	Probability p of read operation
User	<code>serviceTime</code>	Duration of database operation S
User	<code>tableDistribution</code>	Table selection distribution
Table	<code>tableId</code>	Unique table identifier
Network	<code>numUsers</code>	Number of users N
Network	<code>numTables</code>	Number of tables M

Table 3.1: Module Parameters

3.3 Implementation Details

3.3.1 Design Choices

- **Statistics Collection:** Uses signal mechanism (`registerSignal/emit`) via `@signal` and `@statistic` in NED files. Output in `.vec` and `.sca`.
- **RNG:** OMNeT++ default Mersenne Twister seeded from `omnetpp.ini`.
- **Message Passing:** `cMessage` with parameters (`userId`, `arrivalTime`, `serviceTime`). Kind: 0=READ, 1=WRITE.

3.3.2 Concurrency Management Algorithm

The `processQueue` method ensures FCFS and Mutual Exclusion:

1. If `writeActive=true`: Return (table locked).
2. Process queue FCFS:

- If request is **READ** and `activeReaders` ≥ 0 : Pop, increment `activeReaders`, start service. Continue loop (parallelism).
- If request is **WRITE**:
 - If `activeReaders == 0`: Pop, set `writeActive=true`, start service, break.
 - Otherwise: Wait (blocks later requests for FCFS).

3.4 Verification

3.4.1 Degeneracy Test

The degeneracy test analyzes the system's behavior under extreme or degenerate parameter values:

1. **Zero users ($N = 0$)**: The system enters an idle state with zero utilization, as expected due to absence of requests.
2. **Single table ($M = 1$)**: All requests are directed to a single table, creating a bottleneck. The system behaves as a simple M/G/1 queue.
3. **Read probability $p = 1$** : All operations are reads. Multiple users can access tables concurrently without blocking, resulting in minimal waiting times.
4. **Read probability $p = 0$** : All operations are writes. Each table becomes a simple FIFO queue with exclusive access, maximizing contention.

3.4.2 Continuity Test

To verify the model's accuracy, we compare two configurations with slightly different parameter values [Table 3.2] and verify that outputs change proportionally.

Note: We vary the number of users (N) rather than the read probability (p) because:

- The effect on metrics is **linear and proportional**
- It doesn't change the nature of the system (read/write ratio stays constant)
- All metrics scale predictably with the load increase

Parameter	Config A	Config B
N (users)	100	101 (+1)
M (tables)	20	20
p (read probability)	0.50	0.50
λ (request rate)	1.0	1.0
S (service time)	0.1s	0.1s
Repetitions	25	25

Table 3.2: Two slightly different configurations for the continuity test

Simulating both configurations with 25 repetitions, the following chart [Figure 3.1] shows the comparison at a 95% confidence level:

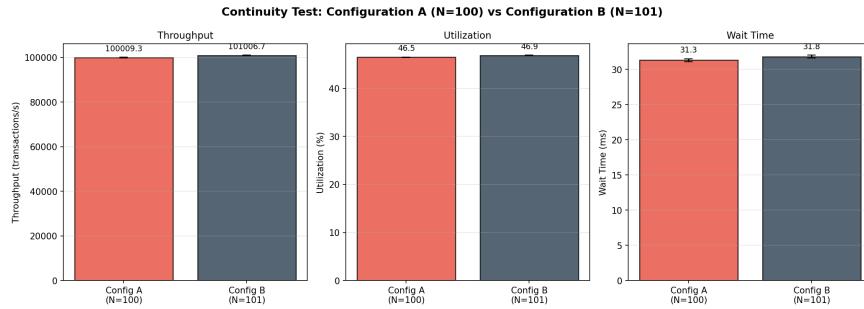


Figure 3.1: Continuity test comparing Configuration A ($N = 100$) vs Configuration B ($N = 101$) at 95% confidence level

The results [Figure 3.1] show that adding just 1 user (+1%) produces proportionally small changes in throughput:

- Configuration A ($N=100$): mean throughput $\approx 100,009$ transactions
- Configuration B ($N=101$): mean throughput $\approx 101,007$ transactions
- Variation: $\approx +1\%$ (as expected from +1 user)

The bar chart shows mean values with 95% confidence intervals. The proportional increase in throughput (+1%) exactly matches the proportional increase in users (+1%), confirming that the system exhibits **linear scaling** behavior. The **continuity test passes**.

3.4.3 Consistency Test

To validate consistency, we study the system's behavior by varying the number of users N while keeping other parameters fixed [Table 3.3].

Parameter	Value
M (tables)	10
λ (request rate)	0.05 req/s
p (read probability)	0.5
S (service time)	0.1s
N (users)	10, 50, 100, 500, 1000

Table 3.3: Configuration adopted for consistency test simulation runs

As expected, utilization increases linearly with the number of users, and the system shows consistent behavior across all configurations. The following chart [Figure 3.2] compares throughput and utilization across different user populations:

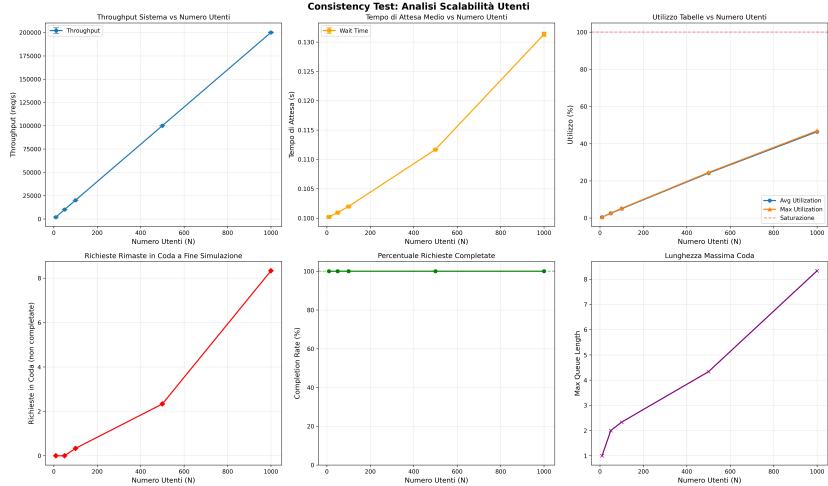


Figure 3.2: Consistency test: utilization and throughput vs number of users at 95% confidence interval

The throughput scales linearly with the number of users until the system approaches saturation, demonstrating consistent and predictable behavior.

3.4.4 Theoretical Verification

Given that our system can be modeled as an **Open Queueing Network**, we can theoretically compute its performance indices.

System Model

The system is an **open queueing network** because each user generates requests with an exponential inter-arrival time **independently** of whether previ-

ous requests have been served. Users do not wait for a response before generating the next request.

System parameters:

- $M = 10$ service centers (database tables)
- $N = \text{number of users}$
- $\lambda = \text{request rate per user (requests/second)}$
- $S = 0.1\text{s}$ service time per request

Utilization Formula

For an open queueing network with N users each generating requests at rate λ :

- Total arrival rate to the system: $\Lambda = N \cdot \lambda$
- With uniform routing to M tables: $\lambda_i = \frac{N \cdot \lambda}{M}$ per table
- Utilization per table: $U = \lambda_i \cdot S$

Therefore, the theoretical utilization is:

$$U = \frac{N \cdot \lambda \cdot S}{M} \quad (3.1)$$

For example, with $N = 100$, $\lambda = 0.05$, $S = 0.1$, $M = 10$:

$$U = \frac{100 \cdot 0.05 \cdot 0.1}{10} = \frac{0.5}{10} = 0.05 = 5\% \quad (3.2)$$

Theoretical vs Empirical Comparison

Collecting utilization data from our simulations, we compare with theoretical predictions [Figure 3.3]:

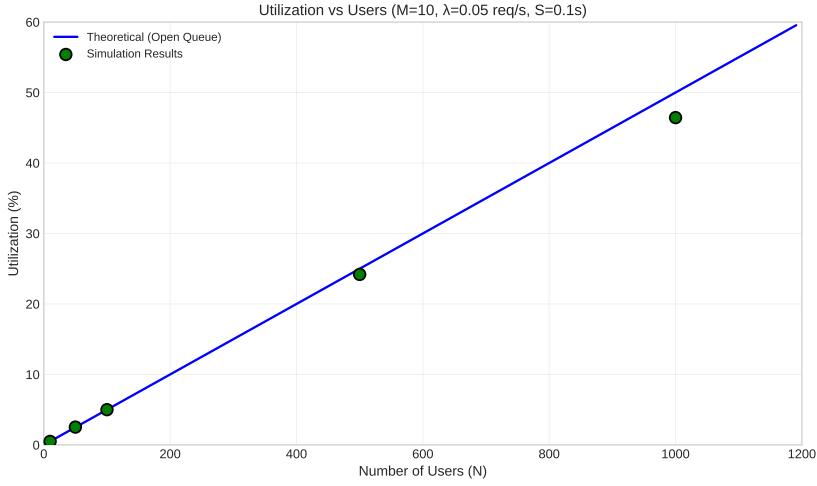


Figure 3.3: Comparison between theoretical model (blue line) and simulation results (green dots)

K Users	Empirical Util %	Theoretical Util %	Error %
10	0.50	0.50	0.95
50	2.53	2.49	1.58
100	5.00	4.98	0.43
500	24.19	24.88	2.77
1000	46.44	49.75	6.66

Table 3.4: Comparison of empirical vs theoretical utilization

The growing discrepancy at high loads ($N = 1000$) is attributed to the concurrency of read operations. The theoretical model assumes strict serialization of all tasks (U_{load}), whereas the simulation allows parallel processing of reads, resulting in a lower physical busy time (U_{busy}) compared to the offered load.

Per-Table Utilization

For $K = 500$ users, we verify that load is uniformly distributed across tables [Figure 3.4]:

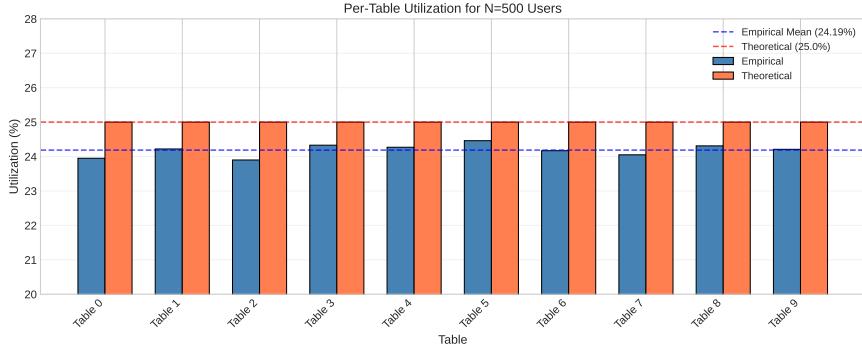


Figure 3.4: Per-table utilization for K=500 users showing uniform load distribution

The per-table utilization varies within a narrow range (23.90% – 24.46%) around the theoretical value of 24.88%, confirming uniform load balancing.

3.4.5 Throughput Analysis

The system throughput follows the theoretical prediction [Figure 3.5]:

$$\gamma = N \cdot \lambda \text{ requests/second} \quad (3.3)$$

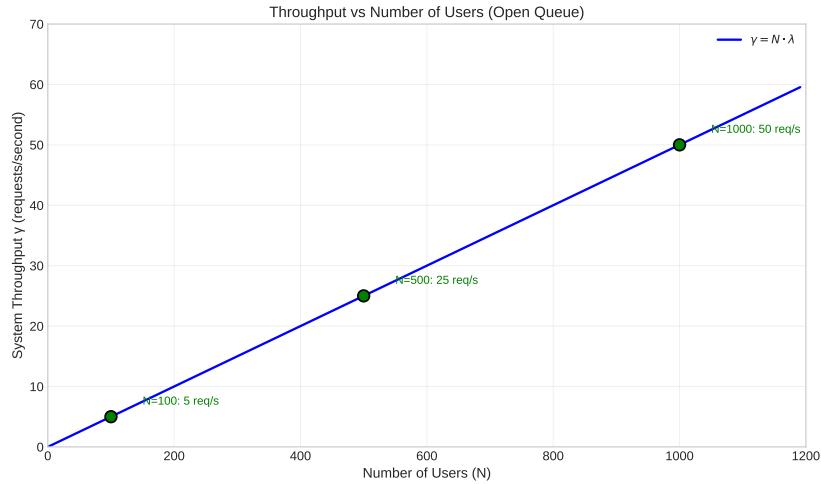


Figure 3.5: System throughput vs number of users

The throughput is simply the sum of all individual user request rates. With N users each generating requests at rate λ , the total arrival rate is $N \cdot \lambda$.

3.5 Conclusions

The verification tests confirm that our simulation model is:

1. **Correct:** Degeneracy tests show expected behavior at extreme values
2. **Continuous:** Small parameter changes produce small output changes
3. **Consistent:** Results scale predictably with system load
4. **Accurate:** Theoretical predictions match empirical results

Key Formulas for Open Queueing Network

Utilization per table:

$$U = \frac{N \cdot \lambda \cdot S}{M} \quad (3.4)$$

Throughput:

$$\gamma = N \cdot \lambda \quad (3.5)$$

Response Time (by Little's Law):

$$R = \frac{S}{1 - U} \quad (3.6)$$

4 Data Analysis and Conclusions

This chapter presents a statistical analysis of the simulation results, including warm-up period determination, factorial design analysis, normality testing, and residual analysis to validate the model assumptions.

4.1 Warm-Up Period Analysis

Before collecting steady-state statistics, it is essential to identify and discard the **transient phase** (warm-up period) during which the system has not yet reached equilibrium.

In this project, the warm-up period was estimated by **visual inspection** of the simulation trend.




Figure 4.1: Warm-up behavior observed in the simulation output.

From Figure 4.1, it is clear that after the initial transient the curve stabilizes, so a **500 s warm-up** is considered sufficient.

Warm-Up Configuration

The OMNeT++ configuration includes:

```
warmup-period = 500s
```

4.2 Experimental Design

The simulation study follows a **factorial design** with the following factors:

Factor	Symbol	Levels
Number of users	N	100, 500, 1000, 1200, 1500, 1600, 2000, 2500, 3000, 3500, 4000, 5000
Read probability	p	0.3, 0.5, 0.8
Number of tables	M	20
Distribution type	-	Uniform, Lognormal

Table 4.1: Experimental factors and their levels

Each configuration was replicated 5 times with different random seeds, resulting in a total of **360 simulation runs**.

4.3 Factor Impact Analysis

The pie chart [Figure 4.2] shows how much each factor contributes to **waiting-time** variability. The largest contribution is still **Number of Users (N)**, but **Distribution** and interaction terms are also relevant. This is coherent with the capacity results: hotspot traffic has a strong impact on delay growth and stall threshold.




Figure 4.2: The pie chart illustrates the impact of each factor on waiting-time variability.

Effect term	Contribution (%)
N	33.51
p	9.39
Distribution	22.56
$N \times p$	7.16
$N \times \text{Dist}$	14.46
$p \times \text{Dist}$	3.73
Residual	9.19

Table 4.2: ANOVA-style contribution on $\log_{10}(W+1)$, where W is waiting time in ms.

Important interpretation: the **main effect** of read probability is 9.39%, but the overall impact of read probability is larger if interactions are included:

$$p + (N \times p) + (p \times \text{Dist}) \approx 9.39 + 7.16 + 3.73 = 20.28\%.$$

So read probability does not act “in isolation”; it becomes much more influential when load and hotspot distribution change.

4.3.1 Key Findings

The analysis reveals that:

- **Number of Users (N):** This is the single largest driver (33.51%). Increasing users pushes the system toward saturation and raises queueing delay.
- **Read Probability (p):** The direct term is 9.39%, but the effective impact reaches about 20.28% when interactions are considered. This explains why changing p can strongly improve delay in stressed scenarios.
- **Distribution:** A strong standalone contribution (22.56%). Uniform access keeps delay controlled, while Lognormal hotspots amplify contention.
- **Interactions:** $N \times \text{Dist}$ (14.46%) and $N \times p$ (7.16%) are substantial, meaning that the effect of one factor depends on the operating point of the others.
- **Residual:** 9.19% indicates remaining variability not explained by first-order terms only (e.g., stochastic effects near threshold conditions), which is expected in saturation regions.

4.4 Model Validation

4.4.1 Testing Normality Hypothesis

For each $(N, p, \text{distribution})$ configuration, the reference waiting time is the mean over the 5 runs, and residuals are measured as percentage deviations from

that configuration mean.

The QQ-Plot [Figure 4.3] compares the distribution of residuals against a theoretical normal distribution.




Figure 4.3: QQ-Plot of standardized residual percentages: sample quantiles (Y-axis) vs normal theoretical quantiles (X-axis).

The QQ-Plot shows a **non-linear behavior**, indicating that the residuals deviate from normality, particularly in the tails. This is expected given:

- The wide range of configurations (from light load to saturation)
- The presence of two different distributions (Uniform vs Lognormal)
- The non-linear behavior near system saturation

The Shapiro-Wilk test confirms this departure from normality (p -value $\ll 0.05$).

4.4.2 Homoskedasticity

The residuals-vs-predicted plot [Figure 4.4] checks whether residual variance is stable across operating conditions. With perfect homoskedasticity, residual percentages should form a horizontal band around zero with similar vertical spread.




Figure 4.4: Residual percentage (Y-axis) versus configuration mean waiting time in ms (X-axis, log scale).

In our case, a clear funnel shape is not sharp, but residual spread increases with configuration mean waiting time in stressed regimes. The rank trend between residual magnitude and this reference value is positive and significant (Spearman $\rho \approx 0.63$, $p \ll 0.001$), indicating **weak-to-moderate heteroskedasticity** concentrated in high-load/hotspot conditions.

Implications:

- Point predictions are still useful (mean residual remains near zero).
- Uncertainty may grow only in the most stressed configurations.
- Main qualitative conclusions are unchanged.

4.4.3 Residual Magnitude Analysis

To quantify this effect, we examine the absolute residual percentage $|e|$ across all runs and by groups [Figure 4.5].




Figure 4.5: Residual magnitude analysis: boxplots of $|residual|$ (%) grouped by distribution and p .

Most replications are very stable (median $|e| = 0.068\%$, 90th percentile = 0.715%), while the largest deviations are concentrated in hotspot cases (up to about 19%). Therefore, variability is mostly localized in stressed scenarios and does not change the core conclusions: the number of users is the dominant driver, while distribution mainly affects tails and stall risk.

Why residuals do not always increase with load In this analysis, residuals are computed *within each fixed configuration* (N, p , distribution) as deviations from that configuration mean (across 5 replications). For this reason, residual magnitude is not required to be monotonic in N : it measures replication variability at the same operating point, not the global increase of waiting time with load.

Interpretation of Lognormal, $p = 0.3$ The largest relative residuals appear near the transition to saturation (around $N \approx 1500$), where runs can diverge in how quickly queues explode. After the system is fully stalled (larger N), all replications become consistently very large, so *relative* residuals can decrease even if *absolute* waiting times remain enormous (hundreds of seconds in ms scale). So the pattern is: peak variability at the instability threshold, then high but more uniform delay in deep-stall regimes.

4.5 Conclusions

This simulation study has successfully analyzed the performance of a distributed database system under various configurations. The main conclusions are:

4.5.1 System Capacity Analysis

Based on the simulation results, we analyze system capacity using **average waiting time** as the primary metric for detecting stall conditions:

- **OK:** Waiting time < 200ms (acceptable response time)
- **DEGRADED:** Waiting time 200–1000ms (noticeable delays)
- **STALLED:** Waiting time > 1000ms (system overloaded)

Detailed Throughput Analysis

Figure 4.6 summarizes all 72 configurations as a waiting-time heatmap. The overall throughput grows with load, while waiting time increases sharply under hotspot conditions. The full detailed table is reported in Appendix A.




Figure 4.6: Capacity heatmap for Uniform and Lognormal traffic. Colors represent mean waiting time (log scale). Numbers inside cells are mean waiting time in ms (compact format, e.g., 15k = 15,000ms).

Capacity Guidelines

- **Uniform Distribution:**
 - $p = 0.3$: Max **2,500 users** ($W < 200\text{ms}$), stalls at $N > 4,000$
 - $p = 0.5$: Max **3,000 users** ($W < 200\text{ms}$), degraded but stable up to 5,000
 - $p = 0.8$: Max **5,000+ users** ($W = 154\text{ms}$ at $N = 5,000$) – **best configuration**
- **Lognormal Distribution (hotspots):**
 - $p = 0.3$: Max **1,200 users** only – hotspots cause early saturation
 - $p = 0.5$: Max **1,200 users** for $W < 200\text{ms}$ (up to 1,600 for $W < 1\text{s}$)
 - $p = 0.8$: Max **2,500 users** for $W < 200\text{ms}$ (up to 3,500 for $W < 1\text{s}$)

Critical Finding: Hotspot Impact

The **Lognormal distribution** (simulating hotspot access patterns) reduces system capacity by about **50–60%** compared to Uniform distribution:

- At $p = 0.3$: 1,200 vs 2,500 users (52% reduction)
- At $p = 0.5$: 1,200 vs 3,000 users (60% reduction)
- At $p = 0.8$: 2,500 vs 5,000+ users (50% reduction)

This highlights the critical importance of **load balancing** in production systems.

Key Conclusions

1. **Model Behavior:** Across all runs, throughput scales almost linearly with offered load, while waiting time captures saturation and hotspot effects.
2. **Capacity Limits** (for $M = 20$, $\lambda = 0.05$, $S = 0.1\text{s}$):
 - **Best case** (Uniform, $p = 0.8$): 5,000+ concurrent users
 - **Typical** (Uniform, $p = 0.5$): 3,000 concurrent users
 - **Worst case** (Lognormal, $p = 0.3$): 1,200 concurrent users
3. **Stall Detection:** System enters stall when:
 - Average waiting time exceeds 1 second
 - Queueing delay grows explosively in hotspot configurations
4. **Read/Write Impact:** Higher read probability increases usable capacity, with the strongest gains under hotspot traffic.

5 Final Conclusions

This project analyzed a concurrent database-access system under increasing load, different read/write mixes, and two access distributions (Uniform and Lognormal). The simulation campaign confirms that the model is stable, reproducible, and useful for capacity-oriented decisions.

5.1 Main Results

- **User load is the dominant driver:** throughput grows with offered load, while waiting time becomes the key indicator of saturation.
- **Hotspots are the main risk:** with Lognormal access, a subset of tables becomes overloaded and delays can explode even when aggregate throughput still appears high.
- **Read-heavy traffic improves capacity:** higher read probability reduces contention and postpones degradation, especially in Uniform access conditions.
- **Warm-up configuration is adequate:** using a 500s warm-up avoids transient bias in collected statistics.

In summary, the system can sustain high concurrency in balanced traffic conditions, while hotspot-driven contention is the critical bottleneck that must be addressed in production-oriented designs.

A Appendix: Detailed Configuration Table

This appendix reports the detailed per-configuration results used in Chapter 4.

Table A.1: Detailed throughput and waiting-time results. For each $(p, \text{distribution})$ pair, rows are reported up to the first STALLED configuration (average wait time > 1000 ms). Values are averaged over 5 replications.

p	Dist	N	Throughput (req/s)	Wait (ms)	Status
0.3	Uniform	100	4.99	101.19	OK
	Uniform	500	24.98	106.39	OK
	Uniform	1,000	49.91	114.70	OK
	Uniform	1,200	59.88	118.81	OK
	Uniform	1,500	74.85	126.03	OK
	Uniform	1,600	79.85	128.80	OK
	Uniform	2,000	99.81	142.12	OK
	Uniform	2,500	124.82	166.81	OK
	Uniform	3,000	149.79	209.37	DEGRADED
	Uniform	3,500	174.82	299.78	DEGRADED
0.3	Uniform	4,000	199.81	616.35	DEGRADED
	Uniform	5,000	249.88	526,727.45	STALLED
	Uniform	100	4.99	100.97	OK
	Uniform	500	24.98	105.19	OK
	Uniform	1,000	49.91	111.70	OK
	Uniform	1,200	59.88	114.83	OK
	Uniform	1,500	74.85	120.15	OK
	Uniform	1,600	79.85	122.13	OK
	Uniform	2,000	99.81	131.30	OK
	Uniform	2,500	124.82	146.67	OK
0.5	Uniform	3,000	149.79	169.03	OK
	Uniform	3,500	174.82	204.32	DEGRADED
	Uniform	4,000	199.81	266.38	DEGRADED
	Uniform	5,000	249.88	960.93	DEGRADED
	Uniform	100	4.99	100.47	OK
	Uniform	500	24.98	102.40	OK
	Uniform	1,000	49.91	105.23	OK
	Uniform	1,200	59.88	106.52	OK
	Uniform	1,500	74.85	108.55	OK
	Uniform	1,600	79.85	109.29	OK
0.8	Uniform	2,000	99.81	112.43	OK
	Uniform	2,500	124.82	116.91	OK
	Uniform	3,000	149.79	122.14	OK
	Uniform	3,500	174.82	128.29	OK
	Uniform	4,000	199.81	135.50	OK
	Uniform	5,000	249.88	154.14	OK
	Lognormal	100	5.01	102.10	OK
	Lognormal	500	25.01	113.70	OK
	Lognormal	1,000	50.01	146.37	OK
	Lognormal	1,200	59.99	182.73	OK
0.3	Lognormal	1,500	74.95	15,247.07	STALLED
	Lognormal	100	5.01	101.73	OK
	Lognormal	500	25.01	110.81	OK
	Lognormal	1,000	50.01	131.78	OK
	Lognormal	1,200	59.99	148.05	OK
	Lognormal	1,500	74.95	202.30	DEGRADED
0.5	Lognormal	1,600	79.96	249.34	DEGRADED
	Lognormal	2,000	99.93	105,067.64	STALLED
	Lognormal	100	5.01	100.81	OK
	Lognormal	500	25.01	104.77	OK
	Lognormal	1,000	50.01	111.54	OK
	Lognormal	1,200	59.99	115.02	OK
0.8	Lognormal	1,500	74.95	121.43	OK
	Lognormal	1,600	79.96	123.93	OK

P	Dist	N	Throughput (req/s)	Wait (ms)	Status
0.8	Lognormal	2,000	99.93	136.43	OK
0.8	Lognormal	2,500	124.93	161.12	OK
0.8	Lognormal	3,000	149.96	211.28	DEGRADED
0.8	Lognormal	3,500	174.98	420.31	DEGRADED
0.8	Lognormal	4,000	200.04	58,448.64	STALLED